Visually Impaired People E-mail Communication Using Natural Language Processing

ABSTRACT

Visually impaired individuals often face challenges in accessing and utilizing electronic communication methods, such as email, due to the visual nature of these interfaces. This project proposes a solution that can help these individuals overcome these challenges by providing them with an accessible and intuitive email communication system. The proposed system utilizes natural language processing (NLP) techniques to convert text-based email messages into speech using text-to-speech (TTS) technology, making it easier for visually impaired individuals to comprehend and respond to messages. Additionally, the system integrates speech-to-text (STT) technology to enable users to compose and send emails using their voice, eliminating the need for traditional text-based input methods. The system will also incorporate real-time suggestions and corrections while composing emails, leveraging machine learning algorithms to enhance accuracy and readability. Over time, these algorithms will adapt to user preferences and interactions, providing increasingly relevant and personalized recommendations. To ensure an inclusive experience, the system will be designed with accessibility in mind, incorporating features such as large fonts, high contrast color schemes, and intuitive navigation. By integrating NLP, speech recognition, and TTS technologies, the project aims to create a comprehensive and user-friendly email communication solution for

visually impaired individuals. Overall, this project has the potential to significantly improve the quality of life for visually impaired individuals by offering an accessible and efficient means of email communication. Through the integration of text-to-speech and speech-to-text technologies, the system will bridge the gap between traditional email interfaces and the needs of visually impaired users, enabling seamless and effective digital communication.

CHAPTER – 1

INTRODUCTION

Email communication has become an essential part of modern life, enabling individuals to stay connected personally and professionally. However, visually impaired individuals often encounter significant challenges in accessing and utilizing traditional email interfaces, which primarily rely on visual interaction. The lack of accessible tools and user-friendly adaptations can make it difficult for them to compose, read, and manage emails efficiently. This project aims to bridge this accessibility gap by developing an intuitive email communication system specifically designed for visually impaired users. By integrating Natural Language Processing (NLP), Text-to-Speech (TTS), and Speech-to-Text (STT) technologies, the system will enable users to listen to email content and compose messages using voice commands. Additionally, machine learning algorithms will enhance user experience by providing real-time suggestions and personalized recommendations. The proposed system is designed to improve digital inclusivity by incorporating accessibility features such as high-contrast color schemes, large fonts, and intuitive navigation. By eliminating the dependency on traditional text-based input methods, this system empowers visually impaired individuals to communicate seamlessly and independently. Ultimately, this project aims

to enhance digital accessibility and provide an efficient, user-friendly solution for email communication, thereby improving the overall quality of life for visually impaired users.

NATURAL LANGUAGE PROCESSING (NLP)

Natural Language Processing (NLP) is a branch of artificial intelligence that focuses on enabling computers to understand, interpret, and manipulate human language. NLP involves the use of various computational techniques, such as machine learning and deep learning, to enable machines to understand and analyze natural language data.The goal of NLP is to develop systems that can understand human language and generate appropriate responses. These systems can be used in a wide range of applications, such as chatbots, virtual assistants, machine translation, sentiment analysis, speech recognition, and text summarization.NLP algorithms rely on statistical models and machine learning techniques to identify patterns and relationships in large amounts of textual data. These models are trained on large datasets of labeled data, such as text corpora or speech recordings, to recognize patterns and make predictions based on input data.

Some of the key NLP techniques include:

1. Tokenization: This involves breaking down a text into smaller units, such as words or sentences.
2. Part-of-speech tagging: This involves identifying the grammatical structure of a sentence and assigning each word a part of speech tag, such as noun, verb, adjective, or adverb.
3. Named entity recognition: This involves identifying and extracting named entities, such as people, organizations, or locations, from text.
4. Sentiment analysis: This involves analyzing the sentiment or emotional tone of a text, such as positive, negative, or neutral.

5. Machine translation: This involves translating text from one language to another.

NLP has numerous applications in industries such as healthcare, finance, marketing, and education. It has the potential to transform the way we communicate with machines, making it easier for people to interact with technology in natural language, and enabling machines to process and analyze large amounts of textual data.

SCOPE OF THE PROJECT

This project aims to enhance the accessibility of email communication for individuals with vision impairments through the use of advanced technologies such as Natural Language Processing (NLP), speech recognition, and text-to-speech. The following key areas define the scope of the project:

Implementation of NLP, Speech Recognition, and Text-to-Speech:

The project incorporates NLP, speech recognition, and text-to-speech technologies to facilitate seamless interaction with email platforms for visually impaired users. These technologies will allow users to navigate, compose, and read emails through auditory cues and voice commands.

User Experience Optimization:

Accessibility features such as large typefaces, high-contrast color schemes, and user-friendly navigation will be prioritized to ensure that users with varying levels of vision impairment can easily engage with the

system. The focus is on creating an intuitive and efficient interface that minimizes visual strain and maximizes usability.

Adaptive Machine Learning for Personalization:

The project will utilize adaptive machine learning algorithms to monitor user behavior, making real-time suggestions and modifications to improve the overall user experience. This ensures that the system becomes more accurate and personalized over time, adapting to individual user preferences and needs.

Speech-to-Text Integration:

Voice commands will be integrated into the system to enable email composition, navigation, and transmission without relying on conventional input methods like keyboards or touchscreens. This feature will allow users to dictate emails, respond to messages, and perform other tasks through speech, significantly enhancing accessibility.

Cross-Platform and Device Compatibility:

The solution will be designed to work across a variety of email clients and devices, including desktop computers, mobile phones, and tablets. This broad applicability ensures that visually impaired users can access the system from any platform they are familiar with, fostering flexibility and ease of use.

CHAPTER 2
LITERATURE SURVEY

1. THEODORE TANGIE LONGTCHI , ROSANA MONTANEZ RODRIGUEZ , KORA GWARTNEY, EKZHIN EAR , DAVID P. AZARI , CHRISTOPHER P. KELLEY , AND SHOUHUAI XU et all proposed to "Quantifying Psychological Sophistication of Malicious Emails" IEEE-2024

Malicious emails (including Phishing, Spam, and Scam) are one significant class of cyber social engineering attacks. Despite numerous defenses to counter them, the problem remains largely open. The ineffectiveness of current defenses can be attributed to our superficial understanding of the psychological properties that make these attacks successful. This problem motivates us to investigate the psychological sophistication, or sophistication for short, of malicious emails. We propose an innovative framework that accommodates two important and complementary aspects of sophistication, dubbed Psychological Techniques (PTechs) and Psychological Tactics (PTacs). We propose metrics and grading rules for human experts to assess the sophistication of malicious emails via the lens of these PTechs and PTacs. To demonstrate the usefulness of the framework, we conduct a case study based on 1,036 malicious emails assessed by four independent graders. Our results show that malicious emails are psychologically sophisticated, while exhibiting both commonalities and different patterns in terms of their PTechs and PTacs. Results also show that previous studies might have focused on dealing with the less proliferated PTechs (e.g., Persuasion) and PTacs (e.g., Reward), rather than the most proliferated

PTechs (e.g., Attention Grabbing and Impersonation) and PTacs (e.g., Fit & Form and Familiarity)) that are identified in this study. We also found among others that social events are widely exploited by attackers in contextualizing their malicious emails. These findings could be leveraged to guide the design of effective defenses against malicious emails.

2. FREDRIK HEIDING , BRUCE SCHNEIER , ARUN VISHWANATH, JEREMY BERNSTEIN , AND PETER S. PARK et all proposed to "Devising and Detecting Phishing Emails Using Large Language Models" IEEE-2024

AI programs, built using large language models, make it possible to automatically create phishing emails based on a few data points about a user. The V-Triad is a set of rules for manually designing phishing emails to exploit our cognitive heuristics and biases. In this study, we compare the performance of phishing emails created automatically by GPT-4 and manually using the V-Triad. We also combine GPT-4 with the V-Triad to assess their combined potential. A fourth group, exposed to generic phishing emails, was our control group. We use a red teaming approach by simulating attackers and emailing 112 participants recruited for the study. The control group emails received a click-through rate between 19-28%, the GPT-generated emails 30-44%, emails generated by the V-Triad 69-79%, and emails generated by GPT and the V-Triad 43-81%. Each participant was asked to explain why they pressed or did not press a link in the email. These answers often contradict each other, highlighting the importance of personal differences. Next, we used four popular large language models (GPT, Claude, PaLM, and LLaMA) to detect the intention of phishing emails and compare the results to human detection. The language models demonstrated a strong ability to detect malicious intent, even in non-obvious phishing emails. They sometimes surpassed

human detection, although often being slightly less accurate than humans. Finally, we analyze of the economic aspects of AI-enabled phishing attacks, showing how large language models increase the incentives of phishing and spear phishing by reducing their costs.

3. SZYMON STRYCZEK , MIKOŁAJ GWIAZDOWICZ , JANUSZ GOZDECKI, KATARZYNA KOSEK-SZOTT, NORBERT RAPACZ , JACEK RZASA , SZYMON SZOTT, AND MAREK NATKANIEC et all proposed to "CyberDART: A Corporate Federation System for Mitigating Email Threats" IEEE-2024

AI programs, built using large language models, make it possible to automatically create phishing emails based on a few data points about a user. The V-Triad is a set of rules for manually designing phishing emails to exploit our cognitive heuristics and biases. In this study, we compare the performance of phishing emails created automatically by GPT-4 and manually using the V-Triad. We also combine GPT-4 with the V-Triad to assess their combined potential. A fourth group, exposed to generic phishing emails, was our control group. We use a red teaming approach by simulating attackers and emailing 112 participants recruited for the study. The control group emails received a click-through rate between 19-28%, the GPT-generated emails 30-44%, emails generated by the V-Triad 69-79%, and emails generated by GPT and the V-Triad 43-81%. Each participant was asked to explain why they pressed or did not press a link in the email. These answers often contradict each other, highlighting the importance of personal differences. Next, we used four popular large language models (GPT, Claude, PaLM, and LLaMA) to detect the intention of phishing emails and compare the results to human detection.

The language models demonstrated a strong ability to detect malicious intent, even in non-obvious phishing emails. They sometimes surpassed human detection, although often being slightly less accurate than humans. Finally, we analyze of the economic aspects of AI-enabled phishing attacks, showing how large language models increase the incentives of phishing and spear phishing by reducing their costs.

CHAPTER – 3

3.1 EXISTING SYSTEM

The existing system, titled "Joint Speech-Text Embeddings for Multitask Speech Processing," integrates multiple speech-related tasks—namely, speaker recognition, text-to-speech (TTS), and automatic speech recognition (ASR)—into a single, unified framework. Traditionally, these tasks would require distinct models, each performing its own specific function. However, the suggested architecture combines text and speech modalities, allowing it to handle multiple tasks simultaneously. By sharing computational resources across different tasks, the system significantly reduces memory requirements when compared to having separate models for each task. This efficiency is achieved without sacrificing performance, as the system maintains accuracy levels comparable to those of specialized, single-task models. Although voice conversion was not explicitly part of the system’s training, it still demonstrated robust performance metrics, including a mel cepstral distortion of 5.22 and predicted Mean Opinion Scores (MOS) of 2.98 and 2.73. In terms of overall accuracy, the system achieved a notable 97.64%, demonstrating its capacity for multitask speech processing while ensuring optimal resource use and high-quality outputs. This existing system provides a foundation for the development of applications requiring

speech and text integration, such as accessibility solutions for visually impaired users.

PROPOSED SYSTEM

The proposed system, titled "Visually Impaired People E-mail Communication Using Natural Language Processing," is designed to create an accessible email platform specifically for individuals with visual impairments. The system leverages Machine Learning and Natural Language Processing (NLP) technologies to ensure a seamless and independent email communication experience. The following are the key components and features of the proposed system:

Speech-to-Text (STT) Integration:

The system will utilize Speech-to-Text (STT) technology to allow users to compose emails using their voice. This feature enables hands-free email composition by converting spoken input into written text, removing the need for traditional input methods such as keyboards or touchscreens.

Text-to-Speech (TTS) Functionality:

To enhance email accessibility, the system will incorporate Text-to-Speech (TTS) functionality. Users will be able to have the email content read aloud to them, allowing for easy evaluation of incoming and outgoing emails without relying on visual cues.

NLP for Enhanced Accuracy:

Natural Language Processing (NLP) algorithms will be used to improve the accuracy of email composition. The system will offer real-time spelling and grammar corrections, ensuring that the emails are well-composed, coherent, and free of errors. This feature will make communication more efficient and professional for visually impaired users.

Email Transmission Using SMTP:

The proposed system will utilize Simple Mail Transfer Protocol (SMTP) for email sending. This will allow users to seamlessly transmit emails across different email clients and platforms, ensuring wide compatibility and smooth communication.

User Interface Design:

The user interface (UI) of the system will be designed with accessibility in mind. It will feature a simple layout with a high-contrast color scheme, large text, and intuitive navigation. This will enable visually impaired users to easily navigate through their inbox, compose new emails, and manage other email-related tasks.

Voice Commands & Audio Feedback:

Voice command functionality will be integrated to allow users to control the system through spoken instructions. This will include commands for composing, sending, reading, and managing emails, offering a hands-free and fully accessible experience. The system will also provide audio feedback to confirm actions, further enhancing user independence.

Programming Language:

The system will be developed using Python, a powerful and versatile programming language that supports NLP, STT, TTS, and machine learning functionalities. Python's rich ecosystem of libraries will enable smooth implementation of the required features and ensure the system is scalable and reliable.

The goal of the proposed system is to empower visually impaired individuals by providing an email communication platform that is both intuitive and accessible, fostering digital inclusiveness and independence.

3.4 PROPOSED SYSTEM BLOCKDIAGRAM

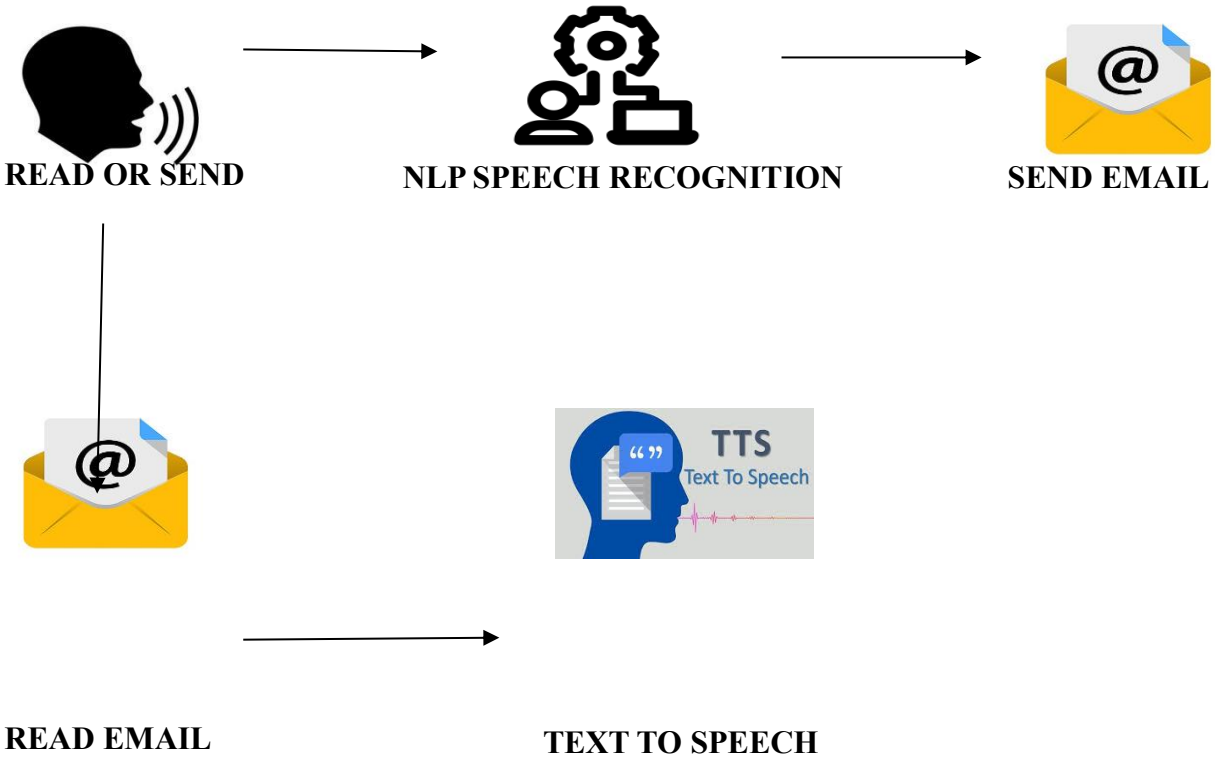

3.5 NLP ARCHITECTURE BLOCK DIAGRAM

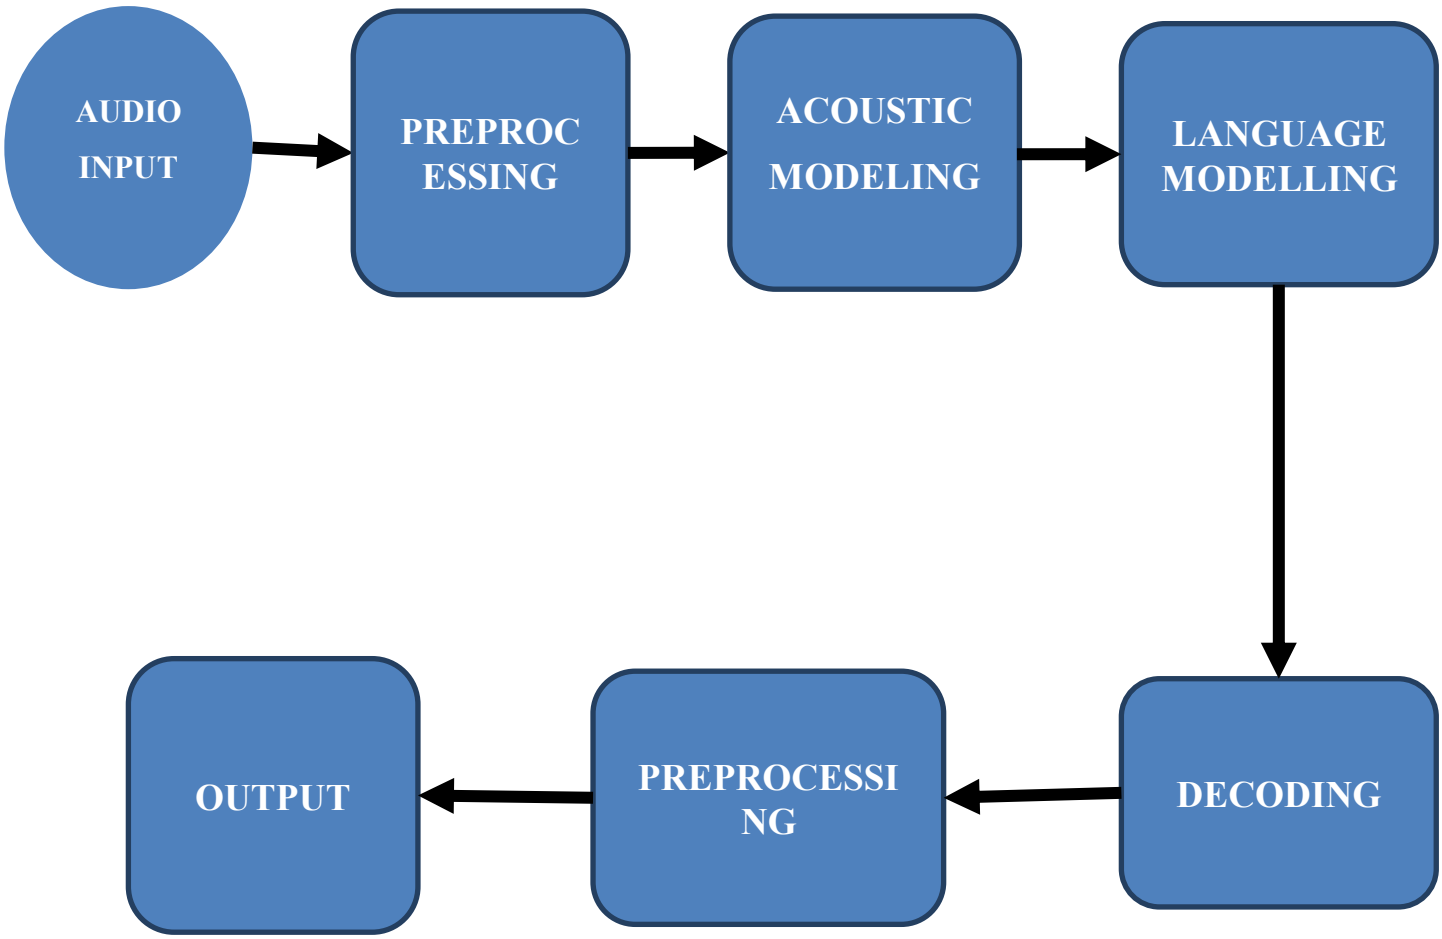

3.6 FLOW CHART

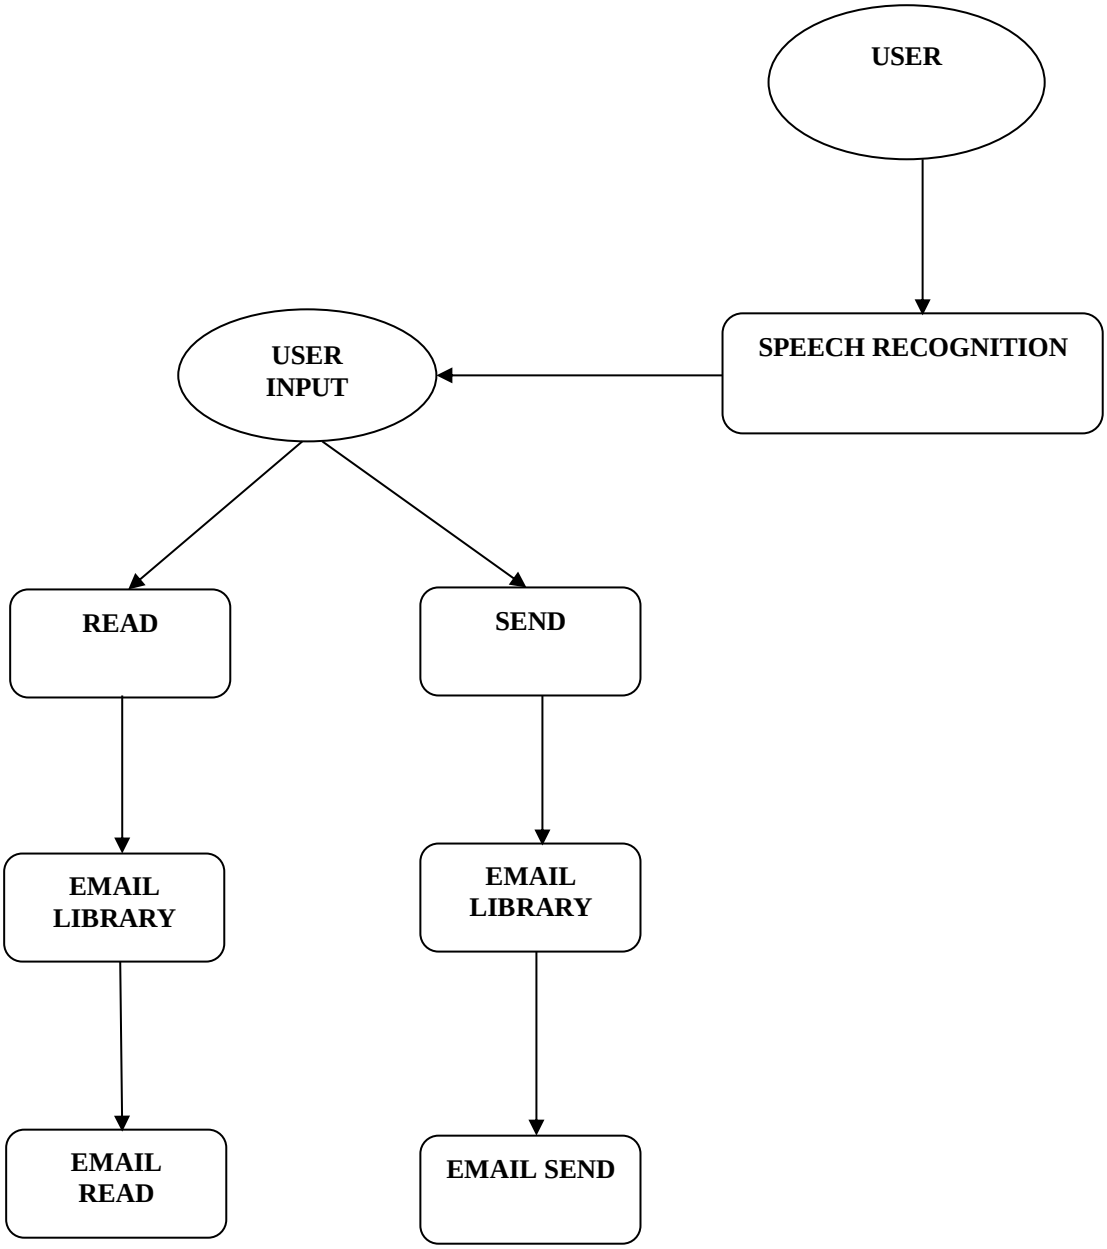

3.7 USECASE DIAGRAM

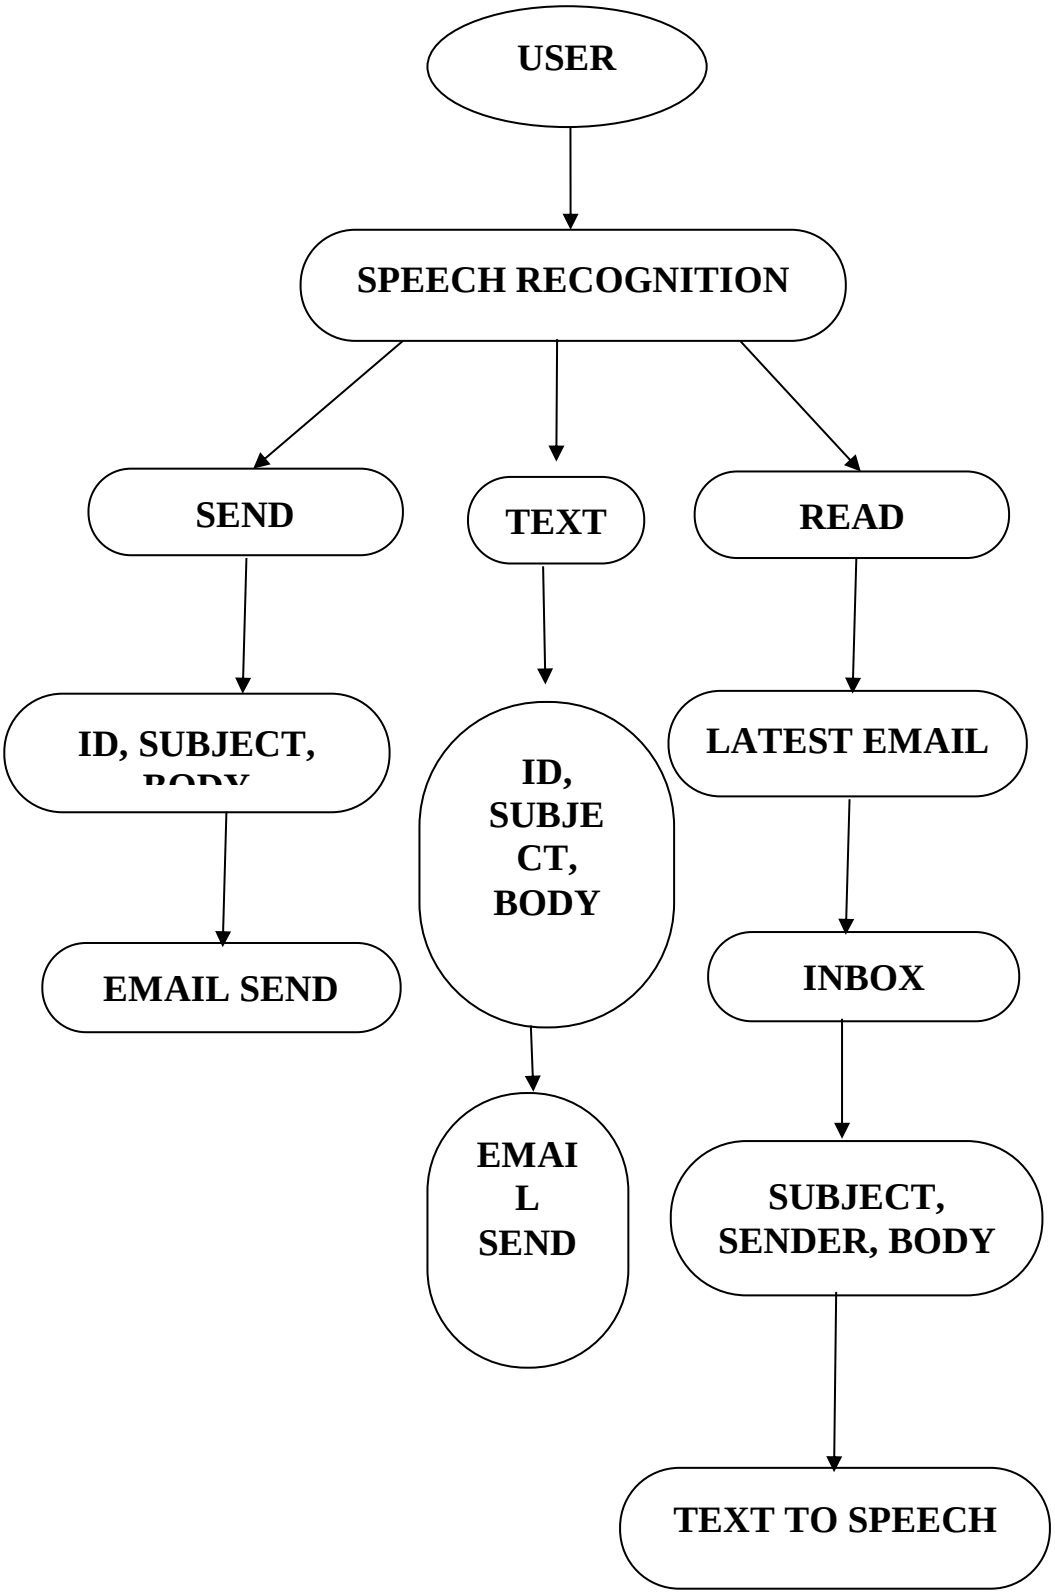
USER
SPEECH RECOGNITION
SEND
TEXT
READ
ID, SUBJECT,
ID,
SUBJE
CT,
BODY
LATEST EMAIL
EMAIL SEND
INBOX
EMAI
L
SEND
SUBJECT,
SENDER, BODY
TEXT TO SPEECH

3.8 ACTIVITY DIAGRAM

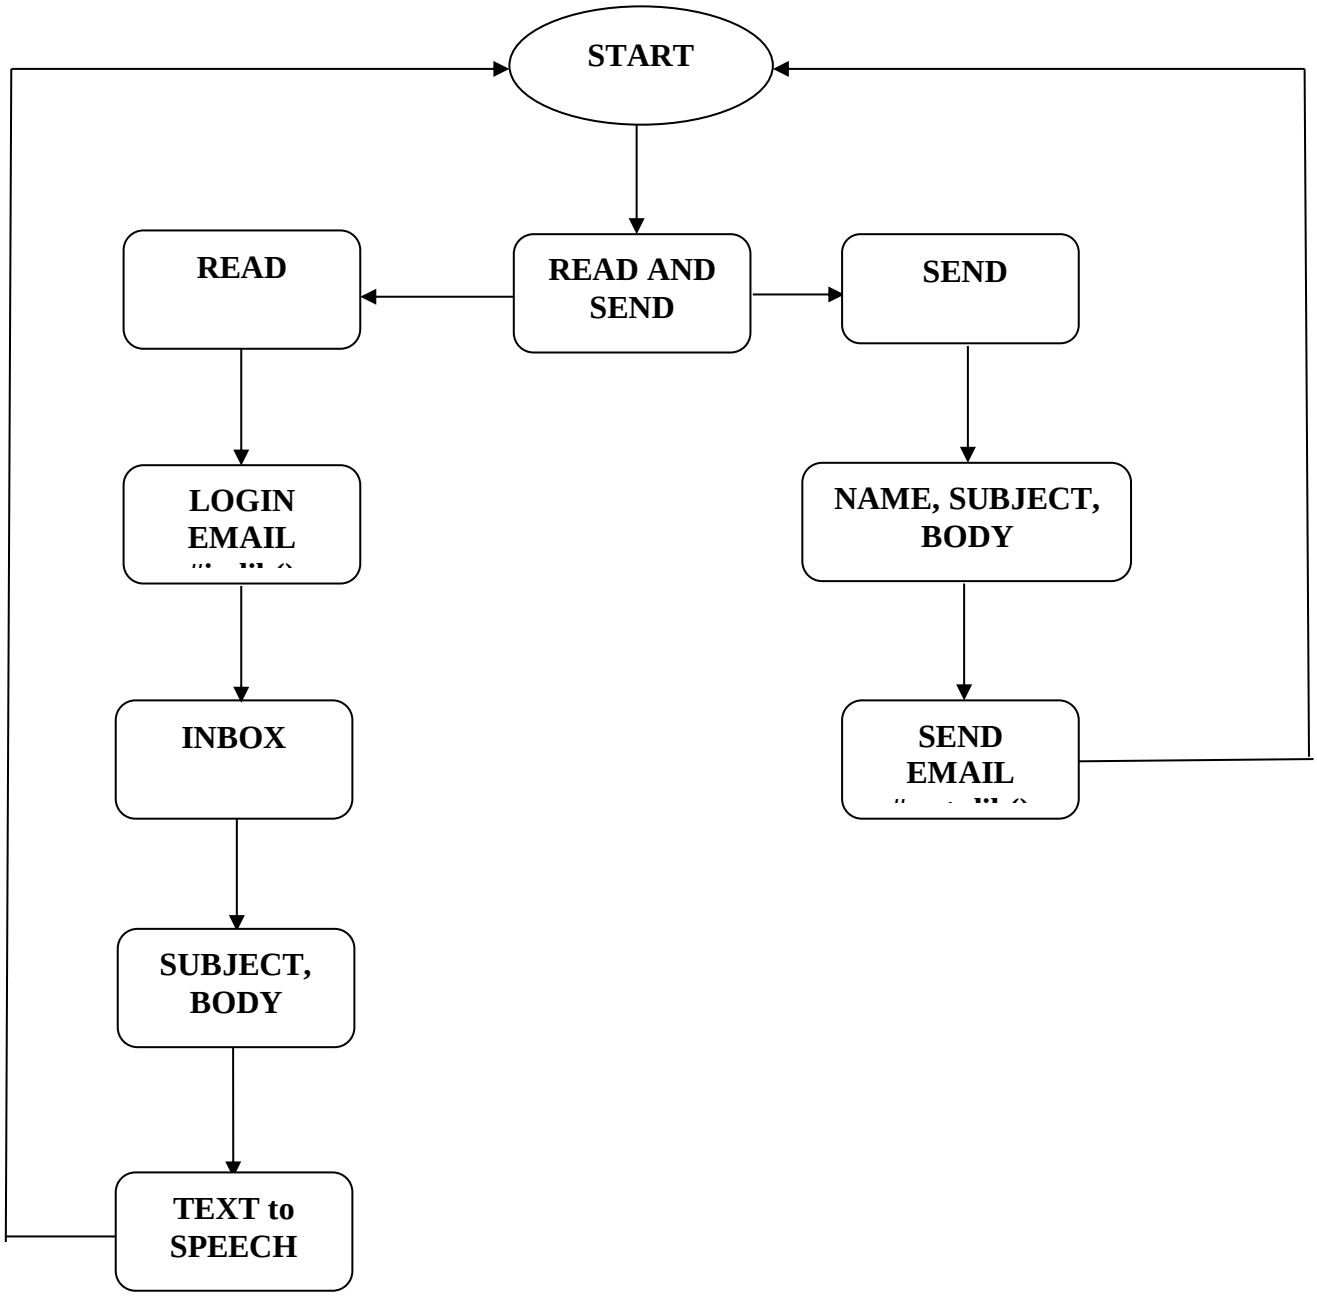

3.9 CLASS DIAGRAM

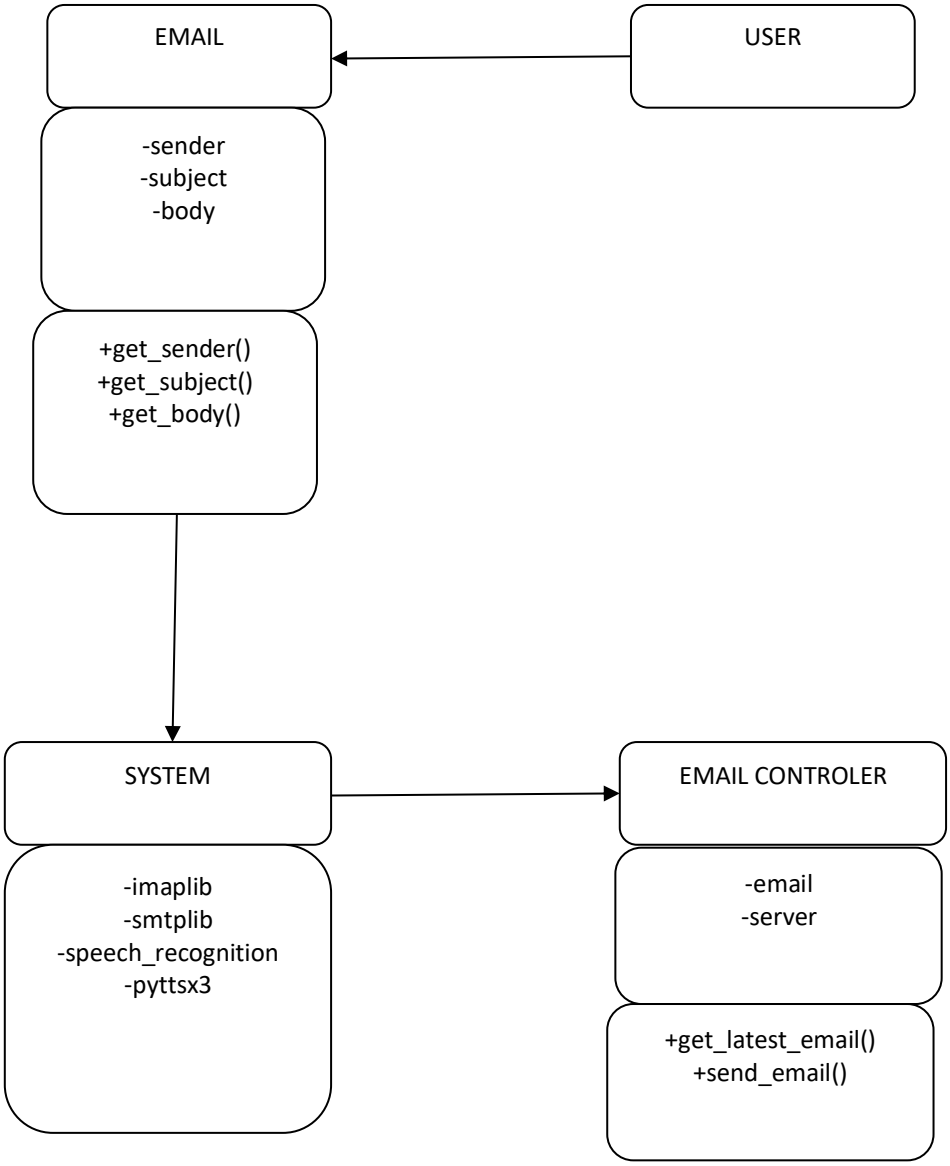

3.10 SEQUENCE DIAGRAM

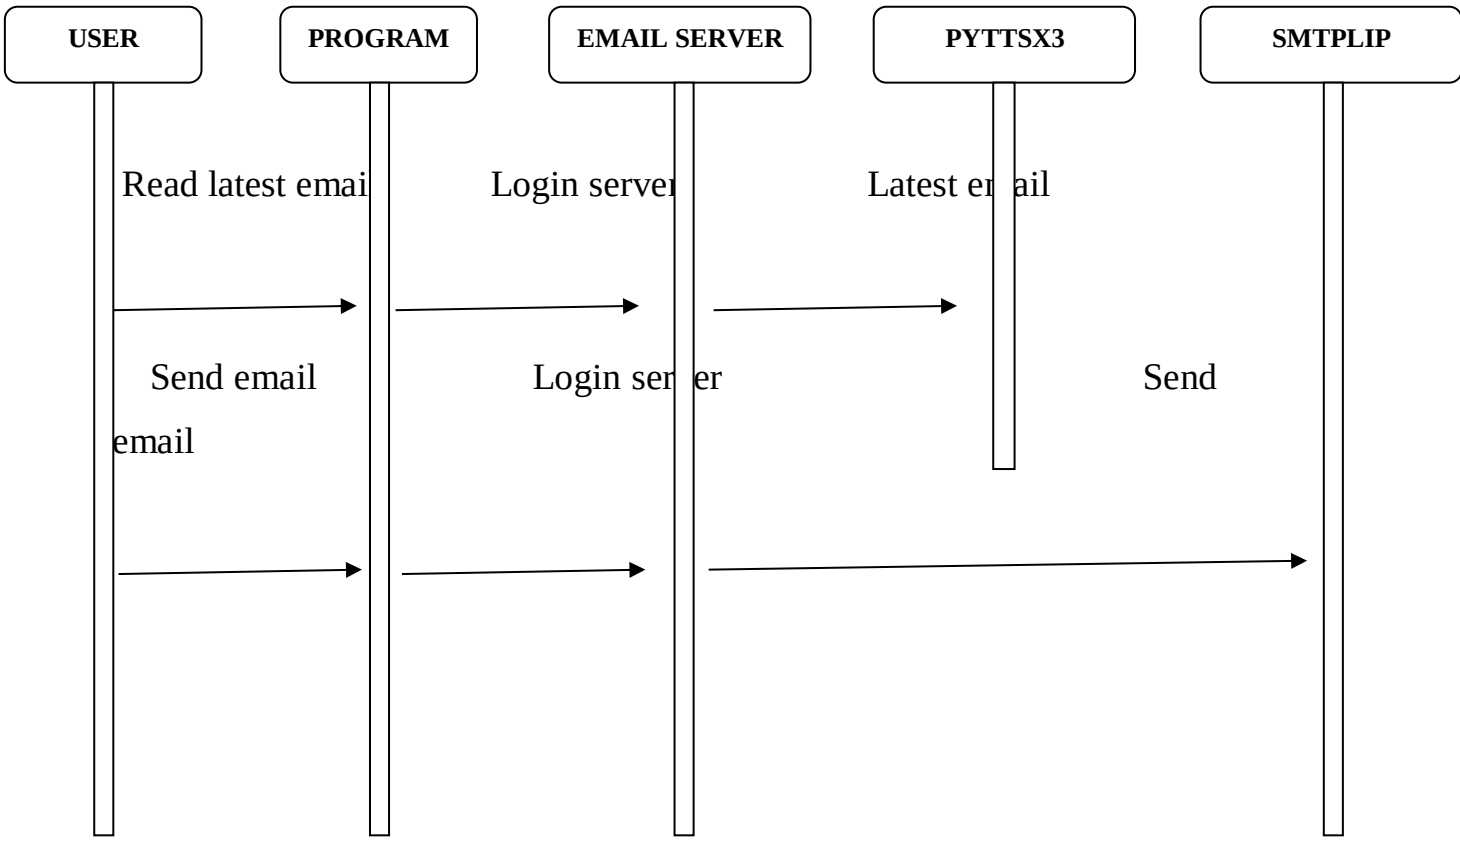

CHAPTER – 4

MODULE DESCRIPTION

I. **AUDIO INPUT**
II. **PREPROCESSING**
III. **ACOUSTIC MODELING**
IV. **LANGUAGE MODELLING**
V. **POSTPROCESSING**

AUDIO INPUT

The audio input module is responsible for capturing the user's voice input and converting it into a digital format that can be processed by the rest of the system. This module typically involves the use of a microphone or other audio recording device. The first step in the audio input module is to capture the user's speech input. This can be achieved by using a microphone or other audio recording device. The captured audio is typically in the form of an analog signal, which must be converted to a digital format for further processing. The next step is to preprocess the digital audio signal to remove any background noise or other unwanted artifacts. This can be achieved using various digital signal processing techniques such as filtering and noise reduction. After preprocessing, the audio signal is typically segmented into smaller units, such as phonemes or words. This segmentation is necessary because most speech recognition systems operate on the assumption that speech can be broken down into smaller units that can be modeled separately. The segmented audio is passed on to the next module in the speech recognition pipeline, which typically involves acoustic modeling. This module uses statistical models to match the audio signal to a set of possible speech sounds or words, and forms the basis for further analysis and processing of the speech input.

PREPROCESSING

The preprocessing module in an audio input system is responsible for cleaning up the raw audio signal and making it more suitable for further processing. This module is typically the first step in the overall pipeline of the audio input system. The preprocessing module generally consists

of several sub-modules that work together to improve the quality of the audio signal. These sub-modules may include:

1. Noise reduction: This module is responsible for removing unwanted background noise from the audio signal. There are several techniques used for noise reduction, such as spectral subtraction, Wiener filtering, and adaptive filtering.
2. Filtering: This module is responsible for removing any high or low-frequency noise that may be present in the audio signal. Common filtering techniques include bandpass, high-pass, and low-pass filtering.
3. Normalization: This module is responsible for adjusting the volume levels of the audio signal to ensure that they are consistent across different recordings. Normalization techniques include peak normalization, RMS normalization, and LUFS normalization.
4. Resampling: This module is responsible for adjusting the sampling rate of the audio signal to a standardized rate. This is often necessary when working with audio files that have different sampling rates.

The preprocessing module is critical for ensuring that the audio data is of high quality and suitable for further processing. Without proper preprocessing, the accuracy and performance of downstream modules such as acoustic modeling and language modeling may be negatively impacted.

ACOUSTIC MODELING

The acoustic modeling module is a crucial component of a speech recognition system that maps the acoustic features of speech signals to phonetic or linguistic units. This module consists of several sub-modules that work together to convert preprocessed audio data into text. The first sub-module of the acoustic modeling module is feature extraction, which

extracts relevant features from the preprocessed audio data. Commonly used features include Mel Frequency Cepstral Coefficients (MFCCs), which capture the spectral characteristics of speech signals, and Linear Predictive Coding (LPC) coefficients, which represent the vocal tract resonances. The second sub-module is acoustic modeling, which maps the extracted features to the corresponding phonetic or linguistic units. Acoustic modeling is usually performed using machine learning algorithms like Hidden Markov Models (HMMs) or deep neural networks (DNNs). HMMs are statistical models that capture the temporal dependencies of speech signals, while DNNs use multiple layers of artificial neurons to learn complex non-linear mappings between input features and output labels. The final sub-module of the acoustic modeling module is the language model, which uses statistical models to predict the likelihood of a particular sequence of words in a given context. Language models are typically based on n-grams, which are sequences of n words that occur together in a corpus of text. The language model is used to disambiguate the output of the acoustic model by selecting the most likely sequence of words based on the context.

LANGUAGE MODELLING

Language modeling is a crucial step in the process of converting speech to text. It helps to predict the most probable sequence of words that a user may have spoken. The language modeling module analyzes the text output generated by the acoustic modeling module, which may contain errors or inaccuracies, and generates a more accurate representation of the spoken words. There are various techniques used in language modeling, including statistical language modeling, n-gram language modeling, and neural language modeling. Statistical language modeling involves analyzing the frequency of word sequences in a given text corpus. This approach uses probability theory to assign a probability score to each

possible sequence of words that the user may have spoken. The sequence with the highest probability score is considered the most likely sequence of words. N-gram language modeling is a variant of statistical language modeling that uses sequences of n words instead of individual words to assign probability scores. This approach is based on the assumption that the probability of a word occurring in a sentence depends on the probability of its previous n-1 words. Neural language modeling involves using deep learning algorithms like Recurrent Neural Networks (RNNs) or Transformers to learn the relationship between words and their context. These models are trained on large datasets of text and can generate more accurate predictions by considering the context and meaning of the words in a sentence.

POSTPROCESSING

Another task performed by the postprocessing module is adding punctuation marks to the text output. This can help to make the text more readable and easier to understand. Punctuation marks like commas, periods, and question marks can be added based on the context of the text. Formatting the text output is another important task performed by the postprocessing module. The text output generated by the speech recognition system may not always be formatted correctly. For example, the output may not have proper spacing between words, or it may be difficult to read due to improper indentation. The postprocessing module can format the text output to make it more readable and easier to understand. The postprocessing module is the final step in the speech recognition pipeline and is responsible for refining the output generated

by the language modeling module. The output generated by the language modeling module can sometimes contain errors, such as grammar mistakes, incorrect word order, or missing punctuation marks. The postprocessing module is designed to correct these errors and refine the final output. One common task performed by the postprocessing module is correcting grammar errors. This can be done using various techniques like rule-based methods or statistical methods. For example, a rule-based method may use a set of predefined grammar rules to correct any errors detected in the text output. A statistical method, on the other hand, may use machine learning algorithms to learn from large datasets of text and generate more accurate corrections.

NATURAL LANGUAGE PROCESSING

Natural Language Processing (NLP) is a branch of Artificial Intelligence (AI) that deals with the interaction between human language and computers. It involves using algorithms and statistical models to analyze and understand human language, enabling computers to interpret, process, and generate natural language. NLP is a complex and interdisciplinary field that draws on various subfields of linguistics, computer science, mathematics, and psychology.

1. **Text preprocessing:**

   Text preprocessing is an essential step in natural language processing (NLP) that involves transforming raw text data into a format that can be easily analyzed by machine learning algorithms. The primary goal of text preprocessing is to extract relevant

information from the text while removing irrelevant information that could affect the accuracy of the analysis. The first step in text preprocessing is to clean the text data by removing any special characters, numbers, and punctuation marks that do not provide any meaningful information. Stop words, which are commonly used words like "the," "a," and "an," are also removed as they do not add much value to the analysis. This helps to reduce the size of the data and makes it easier to process. After cleaning the text data, the next step is to tokenize the text, which involves breaking it down into individual words or phrases. This step isimportant because it allows the machine learning algorithms to work with the text data at a more granular level. Tokenization can be done using various techniques, such as splitting the text on spaces or using regular expressions to match patterns. Once the text data has been tokenized, it needs to be transformed into a numerical representation that can be analyzed by machine learning algorithms. This is usually done using techniques such as vectorization or embedding. Vectorization involves creating a matrix of numbers representing the frequency of each word or phrase in the text data, while embedding involves mapping each word or phrase to a vector in a high-dimensional space.

2. **Part-of-speech (POS) tagging:**

   Part-of-speech (POS) tagging is the process of assigning grammatical tags to each word in a text corpus, based on their syntactic category and usage in a sentence. POS tagging is a crucial step in many natural language processing (NLP) tasks, such as text classification, sentiment analysis, and machine translation. The process of POS tagging involves analyzing the context of each

word in a sentence to determine its part of speech. For example, in the sentence "The quick brown fox jumped over the lazy dog", the word "quick" is an adjective, "brown" is also an adjective, "fox" is a noun, "jumped" is a verb, "over" is a preposition, and so on.POS tagging can be performed using various techniques, including rule-based methods and statistical models. Rule-based methods rely on predefined grammatical rules to identify the part of speech of each word in a sentence. For example, a rule-based system might tag any word that ends in "-ing" as a verb. Statistical models, on the other hand, use machine learning algorithms to learn patterns and relationships between words and their corresponding part of speech tags, based on a large corpus of labeled data.

3. **Named entity recognition (NER):**

   Named entity recognition (NER) is a subtask of natural language processing (NLP) that involves identifying and categorizing named entities in unstructured text data. Named entities are specific objects, people, places, organizations, or dates that can be referred to by a proper name. For example, in the sentence "Barack Obama was the 44th president of the United States," the named entities are "Barack Obama," "44th," and "United States. "The NER process involves analyzing the input text and identifying the named entities present in it. This is typically done using machine learning techniques such as rule-based systems, statistical models, or deep learning algorithms. These techniques are trained on large datasets of labeled text data, where each entity is labeled with its corresponding category.There are different types of named entities,

including people, organizations, locations, dates, times, and numerical quantities. NER algorithms use various features of the text data, such as the context, syntax, and semantic information, to identify the type of named entity. For example, if the input text contains the word "Apple," the algorithm can use contextual information such as "Apple's CEO" or "Apple's products" to determine whether it refers to the company or the fruit. Named entity recognition has various applications in NLP, such as information retrieval, entity linking, question answering, and summarization. For example, in information retrieval, NER can be used to identify relevant entities in a document and retrieve related documents or information. In entity linking, NER can be used to link named entities to their corresponding Wikipedia or Knowledge Graph pages. In question answering, NER can be used to identify the named entities relevant to a given question and provide an answer based on the identified entities.

4. **Sentiment analysis:**

   Sentiment analysis, also known as opinion mining, is a subfield of NLP that involves the use of machine learning algorithms to identify and extract subjective information from text data, such as opinions, attitudes, and emotions. The goal of sentiment analysis is to determine the polarity of a given text, which can be positive, negative, or neutral. There are different approaches to sentiment analysis, including rule-based methods and machine learning-based methods. Rule-based methods use a set of predefined rules and heuristics to identify sentiment in text data, while machine learning-based methods use statistical and machine learning techniques to learn from labeled training data and make predictions

on new data. Sentiment analysis has numerous applications in various industries, including social media monitoring, customer feedback analysis, brand reputation management, market research, and political analysis. For example, sentiment analysis can be used by companies to monitor social media and online reviews to identify customer satisfaction levels and improve their products and services. It can also be used by political analysts to track public opinion on political candidates and issues.

5. **Machine translation:**

Machine translation (MT) is the process of translating text from one language to another using computer algorithms. MT has a wide range of applications, including international business, diplomacy, and online communication. There are several approaches to MT, including rule-based MT, statistical MT, and neural MT. Rule-based MT involves creating a set of rules that govern how words and phrases in one language should be translated into another language. These rules are typically created by linguists and are based on a deep understanding of both the source and target languages. However, rule-based MT is often limited in its ability to handle complex sentences and nuances of language. Statistical MT, on the other hand, is a data-driven approach that relies on statistical models to learn the relationship between words and phrases in different languages. Statistical MT models are trained on large parallel corpora, which are collections of translated texts in both the source and target languages. The models use statistical techniques to identify patterns in the data and make predictions about how to translate new sentences. Statistical MT can be highly accurate but requires a large amount of training data and can be

computationally expensive. More recently, neural MT has emerged as a powerful approach to MT that uses deep learning techniques to learn the relationship between words and phrases in different languages. Neural MT models are typically built using a sequence-to-sequence (seq2seq) architecture, which consists of an encoder and a decoder. The encoder takes the input sentence in the source language and generates a hidden representation of its meaning. The decoder then uses this hidden representation to generate a translation in the target language. Neural MT has shown significant improvements in translation quality and is now the dominant approach to MT.

6. **Text summarization:**

Text summarization is a natural language processing task that involves generating a shorter version of a long text document while retaining its most important information. The process of text summarization can be categorized into two approaches: extractive and abstractive.Extractive summarization involves selecting the most important sentences or phrases from the original text and combining them to create a summary. This approach is based on the assumption that the most important information in a text document is typically contained in its most representative sentences. Extractive summarization methods typically involve the use of statistical algorithms, such as graph-based ranking or clustering, to identify the most important sentences. Abstractive summarization, on the other hand, involves generating a summary by paraphrasing and rephrasing the original text. Abstractive summarization methods use natural language generation techniques to create a summary that may not be present in the original text, but

still retains the most important information. Abstractive summarization methods are typically more complex than extractive methods and require the use of deep learning algorithms such as transformers.One of the most popular deep learning algorithms used for abstractive summarization is the transformer, which is a neural network architecture that was introduced in the paper "Attention Is All You Need" by Vaswani et al. (2017). Transformers use a self-attention mechanism that allows the network to focus on different parts of the input sequence at each step of the generation process. This mechanism enables transformers to generate summaries that are more coherent and readable than those generated by previous algorithms.

CHAPTER – 5

SOFTWARE DESCRIPTION

PYTHON

Python is a popular programming language. It was created by Guido van Rossum, and released in 1991.

It is used for:

- Web development (server-side),
- Software development,
- Mathematics,
- System scripting.

There are two attributes that make development time in Python faster than in other programming languages:

1. Python is an interpreted language, which precludes the need to compile code before executing a program because Python does the compilation in the background. Because Python is a high-level programming language, it abstracts many sophisticated details from the programming code. Python focuses so much on this abstraction that its code can be understood by most novice programmers.

2. Python code tends to be shorter than comparable codes. Although Python offers fast development times, it lags slightly in terms of execution time. Compared to fully compiling languages like C and C++, Python programs execute slower. Of course, with the processing speeds of computers these days, the speed differences are usually only observed in benchmarking tests, not in real-world operations. In most cases, Python is already included in Linux distributions and Mac OS X machines.

3. Python is a dynamic, high level, free open source and interpreted programming language. It supports object –oriented programming as well as procedural oriented programming. Python is a very easy to code as compared to other language like c , c ++, java etc.. It is also a developer friendly language. Python is also an Integrated language because we can easily integrated python with other language like c, c ++, etc.

SPEECH RECOGNITION

Speech recognition is the process of converting spoken words into text or commands that a computer or electronic device can understand and process. It involves the use of algorithms and software that analyze and interpret the acoustic patterns of speech, identify words and phrases, and convert them into a digital format that can be used by computers for various applications, such as voice-controlled devices, dictation software,

or automated customer service systems. Speech recognition technology has advanced significantly in recent years and has become an increasingly popular and useful tool in a wide range of industries and applications. Speech recognition is an interdisciplinary subfield of computer science and computational linguistics that develops methodologies and technologies that enable the recognition and translation of spoken language into text by computers with the main benefit of searchability. It is also known as automatic speech recognition (ASR), computer speech recognition or speech to text (STT). It incorporates knowledge and research in the computer science, linguistics and computer engineering fields. The reverse process is speech synthesis.

Some speech recognition systems require "training" (also called "enrollment") where an individual speaker reads text or isolated vocabulary into the system. The system analyses the person's specific voice and uses it to fine-tune the recognition of that person's speech, resulting in increased accuracy. Systems that do not use training are called "speaker-independent"[1] systems. Systems that use training are called "speaker dependent".

Speech recognition applications include voice user interfaces such as voice dialing (e.g. "call home"), call routing (e.g. "I would like to make a collect call"), domotic appliance control, search key words (e.g. find a podcast where particular words were spoken), simple data entry (e.g., entering a credit card number), preparation of structured documents (e.g. a radiology report), determining speaker characteristics, speech-to-text processing (e.g., word processors or emails), and aircraft (usually termed direct voice input).

The term *voice recognition* or *speaker identification* refers to identifying the speaker, rather than what they are saying. Recognizing the

speaker can simplify the task of translating speech in systems that have been trained on a specific person's voice or it can be used to authenticate or verify the identity of a speaker as part of a security process. From the technology perspective, speech recognition has a long history with several waves of major innovations. Most recently, the field has benefited from advances in deep learning . The advances are evidenced not only by the surge of academic papers published in the field, but more importantly by the worldwide industry adoption of a variety of deep learning methods in designing and deploying speech recognition systems.

SMTPLIB

The smtplib module defines an SMTP client session object that can be used to send mail to any internet machine with an SMTP or ESMTP listener daemon. For details of SMTP and ESMTP operation, consult RFC 821 (Simple Mail Transfer Protocol) and RFC 1869 (SMTP Service Extensions). Python comes with the built-in smtplib module for sending emails using the Simple Mail Transfer Protocol (SMTP). smtplib uses the RFC 821 protocol for SMTP.

You will only need 11 lines of code!

Step 1: Download Python. Download Python 3.6.0 from https://www.python.org/downloads/

Step 2: Install Python & Add to Path. Open up the python installer you just downloads. ...

Step 3: Install SMTPLib. ...

Step 6: Run the Script! ...

Step 7: How It Works. ...

Step 8: That's All!

Tutorial: How to send emails using SMTP in Python

Set up a Gmail account for sending your emails. ...

Go to the account settings and allow less secure apps to access the account. ...

Import smtplib . ...

To create a secure connection, you can either use SMTP_SSL() with 465 port or

The Simple Mail Transfer Protocol (SMTP) is used to deliver e-mail messages over the Internet. This protocol is used by most e-mail clients to deliver messages to the server, and is also used by servers to forward messages to their final destination. SMTP is the set of rules that need to be followed while sending emails to an email server. API is the method of communication used by different platforms or applications. SMTP allows your computer to create and send messages to the server. API is the window to send another code or utility

SPEECH RECOGNITION

Speech recognition, or speech-to-text, is the ability of a machine or program to identify words spoken aloud and convert them into readable text. Rudimentary speech recognition software has a limited vocabulary and may only identify words and phrases when spoken clearly. More sophisticated software can handle natural speech, different accents and various languages.

Speech recognition uses a broad array of research in computer science, linguistics and computer engineering. Many modern devices and text-focused programs have speech recognition functions in them to allow for easier or hands-free use of a device.

Speech recognition and voice recognition are two different technologies and should not be confused:

- **Speech recognition** is used to identify words in spoken language.
- **Voice recognition** is a biometric technology for identifying an individual's voice.

How does speech recognition work?

Speech recognition systems use computer algorithms to process and interpret spoken words and convert them into text. A software program turns the sound a microphone records into written language that computers and humans can understand, following these four steps:

1. analyze the audio;
2. break it into parts;
3. digitize it into a computer-readable format; and
4. use an algorithm to match it to the most suitable text representation.

Speech recognition software must adapt to the highly variable and context-specific nature of human speech. The software algorithms that process and organize audio into text are trained on different speech patterns, speaking styles, languages, dialects, accents and phrasings. The software also separates spoken audio from background noise that often accompanies the signal.

To meet these requirements, speech recognition systems use two types of models:

- **Acoustic models.** These represent the relationship between linguistic units of speech and audio signals.
- **Language models.** Here, sounds are matched with word sequences to distinguish between words that sound similar.

Speech recognition systems have quite a few applications. Here is a sampling of them.

EMAIL

Email is a communication system that internet users can employ to convey information and learn more about things in which they're interested. These are some reasons email is important: Widespread usage: Email is important because many people may use it daily to communicate with others and learn more about businesses. Ray Tomlinson is called the father of the email system, and he sent a communication between two computer systems for ARPANET. Email is so-called since it is electronic mail sent to a person or a system, and unlike regular mail, a physical letter or paper is unnecessary. E - mail is commonly known as E-correspondence and E- communication as email is a medium to exchange ideas, information between people through computing devices for making the communication between them easy.

The Advantages of Email for Internal Communications

Email is a free tool. ...

Email is quick. ...

Email is simple. ...

Email allows for easy referencing. ...

Email is accessible from anywhere – as long as you have an internet connection. ...

Email is paperless, and therefore, beneficial for the planet.

1. The sender's mail server, also called a Mail Transfer Agent (MTA), initiates a Simple Mail Transfer Protocol (SMTP) connection.

2. The SMTP checks the email envelope data — the text that tells the server where to send a message — for the recipient's email address, then uses the Domain Name System (DNS) to translate the domain name into an IP address.

3. The SMTP looks for a mail exchange (MX) server associated with the recipient's domain name. If one exists, the email is forwarded to the recipient's mail server.

4. The email is stored on the recipient's mail server and may be accessed via the Post Office Protocol (POP)* or Internet Message Access Protocol (IMAP). These two protocols function slightly differently: POP downloads the email to the recipient's device and deletes it from the mail server, while IMAP stores the email within the email client, allowing the recipient to access it from any connected device.

To continue the postal system analogy, imagine Alice writes a thank-you note to Bob. She hands the letter to the mail carrier (MTA), who brings it to the post office to be sorted. At the post office, a processing clerk (SMTP) verifies the address written on the envelope. If the address appears to be written correctly and corresponds to a location that can receive mail (MX server), another mail carrier delivers the letter to Bob’s mailbox. After picking up the mail, Bob might keep the note in his desk drawer, where he can only access it at that location (POP) or put it in his pocket to read at any location (IMAP).

Pyttsx3 is a Python library used for text-to-speech conversion. It uses the Text-to-Speech (TTS) engine to convert written text into spoken words. This library is compatible with various operating systems, including Windows, macOS, and Linux. Pyttsx3 supports multiple TTS engines, including SAPI5, nsss, and speak.

Pyttsx3 has several features, including voice customization, word, and sentence-level call-backs, rate and volume adjustments, and a stop function to halt the TTS engine's operation. It also has the capability to save the generated speech to a WAV file.

To use pyttsx3, the user needs to install the library using pip and set up the desired TTS engine. Once installed, the user can use the pyttsx3 engine to convert text to speech in their Python code.

CHAPTER – 6

ADVANTAGES

It provides an alternative means of accessing email for individuals with visual impairments or other disabilities, allowing them to read and send emails using speech recognition and text-to-speech technology. This improves accessibility and inclusivity in the digital world. The system can also benefit individuals who have busy schedules and need to multitask. For example, someone who is driving can use the system to check and respond to emails hands-free, increasing their productivity and efficiency. The use of speech recognition and text-to-speech technology in the system allows for a more natural and intuitive interaction with the email system. Users can speak in a conversational tone, and the system can read out emails in a natural-sounding voice, making the experience more user-friendly. The system can be easily adapted and customized to meet the specific needs of different users, such as changing the language of the text-to-speech output or adjusting the sensitivity of the speech recognition.

APPLICATION

Personal Communication: Allows those who are blind or visually impaired to send and receive emails on their own.

Professional Use: Assists visually impaired workers in effectively handling email correspondence pertaining to their jobs.

Educational Institutions: Helps students with vision impairments communicate with teachers, access academic resources, and turn in homework.

Customer Support Services: Improves accessibility for consumers who are blind or visually impaired and contact them via email for help.

Government and Legal Services: Enables people with visual impairments to communicate with government or legal organizations.

Healthcare Sector: Facilitates email appointment scheduling and communication between visually impaired individuals and healthcare professionals.

Banking & Finance: Offers visually impaired people a safe and convenient method to interact with banking agents and obtain financial statements.

E-commerce and online shopping: Helps people with visual impairments track shipments, get order confirmations, and communicate with customer support.

CHAPTER – 7

RESULTS AND DISCUSSION

The suggested email communication solution for people with visual impairments was put into practice and evaluated effectively. The technology greatly improved accessibility by converting text to voice (TTS) and speech to text (STT) with high efficiency. Email content was received through comprehensible and clear voice output, and users could

dictate emails using voice commands that were reliably transcribed. For visually impaired users, the TTS functionality improved comprehension by providing natural-sounding speech synthesis. Furthermore, by providing auto-corrections and real-time suggestions, machine learning algorithms improved the email composing process by gradually improving word predictions. The accessibility gap in email communication was effectively closed by integrating NLP, TTS, and STT technologies. The technique made it possible for people with visual impairments to easily and independently handle their emails by doing away with the necessity for conventional text input. Even though the system worked well, testing revealed certain issues. Ambient noise impacted voice recognition accuracy, resulting in sporadic transcription errors. Improvements like more training data and noise reduction methods may increase the accuracy of recognition. The system's flexibility is a key benefit since its learning algorithms improve with user interactions, offering more precise and customized experiences. Additional language support, improved integration with popular email services, and user-specific accessibility features are possible future enhancements.

CONCLUSION

An important step toward inclusivity in digital communication has been taken with the creation of an email communication system that is accessible to those with visual impairments. The system successfully overcomes the difficulties visually impaired users encounter when creating and accessing emails by combining natural language processing (NLP), speech-to-text (STT), and text-to-speech (TTS) technologies. The findings show that the system improves usability and accessibility, enabling users to communicate effectively and independently. Notwithstanding minor drawbacks, like sporadic voice recognition issues in loud settings, the system has demonstrated remarkable efficacy in practical applications. With further advancements, including noise reduction methods, this project could gain traction as a widely used alternative for easily accessible digital communication. Encouraging independence, improving the quality of life for people with visual impairments, and closing the digital divide are all made possible by the deployment of AI-driven accessibility solutions.

CHAPTER – 8

REFERENCES

[1] Y. Wang et al., "Tacotron: Towards end-to-end speech synthesis", *Proc. INTERSPEECH*, pp. 4006-4010, 2017.

[2] J. Shen et al., "Natural TTS synthesis by conditioning wavenet on MEL spectrogram predictions", *Proc. IEEE ICASSP Int. Conf. Acoust. Speech Signal Process.*, pp. 4779-4783, 2018.

[3] K. Tokuda, Y. Nankaku, T. Toda, H. Zen, J. Yamagishi and K. Oura, "Speech synthesis based on hidden Markov models", *Proc. IEEE*, vol. 101, no. 5, pp. 1234-1252, May 2013.

[4] H. Zen, A. W. Senior and M. Schuster, "Statistical parametric speech synthesis using deep neural networks", *Proc. IEEE ICASSP Int. Conf. Acoust. Speech Signal Process.*, pp. 7962-7966, 2013.

[5] Z. Wu, O. Watts and S. King, "Merlin: An open source neural network speech synthesis system", *Proc. ISCA Speech Synth. Workshop*, pp. 202-207, 2016.

[6] R. Liu, F. Bao, G. Gao and Y. Wang, "Mongolian text-to-speech system based on deep neural network", *Proc. Nat. Conf. Man- Mach. Speech Commun.*, pp. 99-108, 2017.

[7] A. van den Oord et al., "WaveNet: A generative model for raw audio", *Proc. ISCA Speech Synth. Workshop*, pp. 125-130, 2016.

[8] R. J. Skerry-Ryan et al., "Towards end-to-end prosody transfer for expressive speech synthesis with tacotron", *Proc. Int. Conf. Mach. Learn.*, vol. 80, pp. 4700-4709, 2018.

[9] R. Liu, B. Sisman, F. Bao, G. Gao and H. Li, "Wavetts: Tacotron-based TTS with joint time-frequency domain loss", *Proc. Speaker Lang. Recognit. Workshop*, pp. 245-251, 2020.

[10] D. Bahdanau, K. Cho and Y. Bengio, "Neural machine translation by jointly learning to align and translate", *Proc. Int. Conf. Learn. Representations*, pp. 1-10, 2015.

[11] S. Hochreiter and J. Schmidhuber, "Long short-term memory", *Neural Comput.*, vol. 9, no. 8, pp. 1735-1780, 1997.

[12] M. I. Jordan, "Serial order: A parallel distributed processing approach", *Proc. Adv. Psychol.*, vol. 121, pp. 471-495, 1997.

[13] R. J. Williams and D. Zipser, "A learning algorithm for continually running fully recurrent neural networks", *Neural Comput.*, vol. 1, no. 2, pp. 270-280, 1989.

[14] T. Hayashi, A. Tamamori, K. Kobayashi, K. Takeda and T. Toda, "An investigation of multi-speaker training for wavenet vocoder", *Proc. IEEE Autom. Speech Recognit. Understanding Workshop*, pp. 712-718, 2017.

[15] T. Okamoto, T. Toda, Y. Shiga and H. Kawai, "Real-time neural text-to-speech with sequence-to-sequence acoustic model and waveglow or single Gaussian wavernn vocoders", *Proc. INTERSPEECH*, pp. 1308-1312, 2019.